

# Data-X HW1 Sp18

February 8, 2018

## 1 Homework 1

In this homework, you will get a chance to do some exercises with Numpy, Pandas, and Matplotlib to show us your understanding with this libraries.

If you have questions, Google! Additionally you can ask your peers questions on Piazza and/or go to Office Hours.

This homework is due **Thursday Feb. 8th, 2018 at 11:59 PM**. Please upload your .ipynb to your private repo on Github. Additionally, submit a pdf on bCourses and in the comment section include a link to your private repo.

This homework is long, please start early!

```
In [19]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
```

### 1.1 NumPy Basics

Create two numpy arrays (a and b). a should be all integers between 10-19 (inclusive), and b should be ten evenly spaced numbers between 1-7. Print the results below.

```
In [20]: a = np.arange(10, 20)
b = np.linspace(1, 7, 10)

print("a: ", a)
print("b: ", b)
```

a: [10 11 12 13 14 15 16 17 18 19]  
b: [ 1. 1.66666667 2.33333333 3. 3.66666667 4.33333333  
5. 5.66666667 6.33333333 7. ]

For a and b above do the follow and print out the results.

1. Square all the elements in both arrays (element-wise).
2. Add both the squared arrays (e.g.  $[1,2] + [3,4] = [4,6]$ ).

3. Sum the elements with even indices of the added array.
4. Take the square root of the added array (element-wise square root).

In [21]: # Variables

```
a_sq = np.square(a)
b_sq = np.square(b)
added_array = a_sq + b_sq
indices = np.array([x for x in range(0, len(added_array), 2)])
```

```
print("1. ", "a: ", a_sq, " - b: ", b_sq)
print("2. ", added_array)
print("3. ", np.sum(added_array[indices]))
# print("4. ", np.sqrt(np.sum(added_array)))
print("4. ", np.sqrt(added_array))
```

```
1.  a:  [100 121 144 169 196 225 256 289 324 361] - b:  [ 1.          2.77777778  5.44444444
18.77777778 25.          32.11111111 40.11111111 49.          ]
2.  [ 101.          123.77777778 149.44444444 178.          209.44444444
243.77777778 281.          321.11111111 364.11111111 410.          ]
3.  1105.0
4.  [ 10.04987562 11.12554618 12.22474721 13.34166406 14.47219556
15.61338457 16.76305461 17.91957341 19.08169571 20.24845673]
```

Append b to a. Reshape the appended array so that it is a 5x4, 2D-array and store the results in a variable called m. Print m.

In [22]: m = np.reshape(np.concatenate((a,b)), (5, 4))

```
print("m: ", m)
```

```
m:  [[ 10.          11.          12.          13.          ]
 [ 14.          15.          16.          17.          ]
 [ 18.          19.          1.          1.66666667]
 [ 2.33333333  3.          3.66666667  4.33333333]
 [ 5.          5.66666667  6.33333333  7.          ]]
```

Extract the second and the third column of the matrix m. Store the resulting 5x2 matrix in a new variable called m2. Print m2.

In [23]: m2 = m[:, [1, 2]]

```
print("m2: ", m2)
```

```
m2:  [[ 11.          12.          ]
 [ 15.          16.          ]
 [ 19.          1.          ]]
```

```
[ 3.          3.66666667]
[ 5.66666667  6.33333333]]
```

Take the dot product of `m2` and `m` store the results in a matrix called `m3`. Print `m3`. Note that dot product of two matrices  $A \cdot B = A^T B$

```
In [24]: m3 = np.transpose(m2).dot(m)
```

```
print("m3: ", m3)
```

```
m3: [[ 697.33333333  748.11111111  437.88888889  482.33333333]
      [ 402.22222222  437.88888889  454.55555556  489.88888889]]
```

Round the `m3` matrix to two decimal points. Store the result in place and print the new `m3`.

```
In [25]: m3 = m3.round(2)
```

```
print("m3: ", m3)
```

```
m3: [[ 697.33  748.11  437.89  482.33]
      [ 402.22  437.89  454.56  489.89]]
```

Sort the `m3` array so that the highest value is at the top left, the next highest value to the right of the highest, and the lowest value is at the bottom right. Print the sorted `m3` array.

```
In [26]: sorted_m3 = np.sort(m3.reshape(8,1)).reshape(2,4)
```

```
print("sorted m3: ", sorted_m3)
```

```
sorted m3: [[ 697.33  748.11  437.89  482.33]
             [ 402.22  437.89  454.56  489.89]]
```

## 1.2 NumPy and Masks

Create an array called `f` where there are 100 equally-spaced values from 0 to `pi`, inclusive. Take the sin of the array `f` (element-wise) and store that in place. Print `f`.

```
In [27]: from numpy import pi
```

```
f = np.sin(np.linspace(0, np.pi, 100))
```

```
print("f: ", f)
```

```
f: [ 0.00000000e+00  3.17279335e-02  6.34239197e-02  9.50560433e-02
 1.26592454e-01  1.58001396e-01  1.89251244e-01  2.20310533e-01
 2.51147987e-01  2.81732557e-01  3.12033446e-01  3.42020143e-01
 3.71662456e-01  4.00930535e-01  4.29794912e-01  4.58226522e-01
 4.86196736e-01  5.13677392e-01  5.40640817e-01  5.67059864e-01
 5.92907929e-01  6.18158986e-01  6.42787610e-01  6.66769001e-01
 6.90079011e-01  7.12694171e-01  7.34591709e-01  7.55749574e-01
 7.76146464e-01  7.95761841e-01  8.14575952e-01  8.32569855e-01
 8.49725430e-01  8.66025404e-01  8.81453363e-01  8.95993774e-01
 9.09631995e-01  9.22354294e-01  9.34147860e-01  9.45000819e-01
 9.54902241e-01  9.63842159e-01  9.71811568e-01  9.78802446e-01
 9.84807753e-01  9.89821442e-01  9.93838464e-01  9.96854776e-01
 9.98867339e-01  9.99874128e-01  9.99874128e-01  9.98867339e-01
 9.96854776e-01  9.93838464e-01  9.89821442e-01  9.84807753e-01
 9.78802446e-01  9.71811568e-01  9.63842159e-01  9.54902241e-01
 9.45000819e-01  9.34147860e-01  9.22354294e-01  9.09631995e-01
 8.95993774e-01  8.81453363e-01  8.66025404e-01  8.49725430e-01
 8.32569855e-01  8.14575952e-01  7.95761841e-01  7.76146464e-01
 7.55749574e-01  7.34591709e-01  7.12694171e-01  6.90079011e-01
 6.66769001e-01  6.42787610e-01  6.18158986e-01  5.92907929e-01
 5.67059864e-01  5.40640817e-01  5.13677392e-01  4.86196736e-01
 4.58226522e-01  4.29794912e-01  4.00930535e-01  3.71662456e-01
 3.42020143e-01  3.12033446e-01  2.81732557e-01  2.51147987e-01
 2.20310533e-01  1.89251244e-01  1.58001396e-01  1.26592454e-01
 9.50560433e-02  6.34239197e-02  3.17279335e-02  1.22464680e-16]
```

Use a 'mask' and print an array that is True when  $f \geq 1/2$  and False when  $f < 1/2$ . Print an array sequence that has only those values where  $f \geq 1/2$ .

```
In [28]: mask = f >= 1/2
         print(mask)
         print()
         print(f[mask])
```

```
[False False False False False False False False False False False False
 False False False False False  True  True  True  True  True  True  True
  True  True  True  True  True  True  True  True  True  True  True  True
  True  True  True  True  True  True  True  True  True  True  True  True
  True  True  True  True  True  True  True  True  True  True  True  True
  True  True  True  True  True  True  True  True  True  True  True False
 False False False False False False False False False False False False
 False False False False]
```

```
[ 0.51367739  0.54064082  0.56705986  0.59290793  0.61815899  0.64278761
 0.666769    0.69007901  0.71269417  0.73459171  0.75574957  0.77614646
 0.79576184  0.81457595  0.83256985  0.84972543  0.8660254   0.88145336]
```

```

0.89599377  0.909632    0.92235429  0.93414786  0.94500082  0.95490224
0.96384216  0.97181157  0.97880245  0.98480775  0.98982144  0.99383846
0.99685478  0.99886734  0.99987413  0.99987413  0.99886734  0.99685478
0.99383846  0.98982144  0.98480775  0.97880245  0.97181157  0.96384216
0.95490224  0.94500082  0.93414786  0.92235429  0.909632    0.89599377
0.88145336  0.8660254   0.84972543  0.83256985  0.81457595  0.79576184
0.77614646  0.75574957  0.73459171  0.71269417  0.69007901  0.666769
0.64278761  0.61815899  0.59290793  0.56705986  0.54064082  0.51367739]

```

### 1.3 NumPy and 2 Variable Prediction

Let  $x$  be the number of miles a person drives per day and  $y$  be the dollars spent on buying car fuel per day.

We have created 2 numpy arrays each of size 100 that represent  $x$  and  $y$ .

$x$  (number of miles) ranges from 1 to 10 with a uniform noise of  $(0, 1/2)$ .

$y$  (money spent in dollars) will be from 1 to 20 with a uniform noise  $(0, 1)$ .

Run the cell below.

```

In [29]: # seed the random number generator with a fixed value
         np.random.seed(500)

         x=np.linspace(1,10,100)+ np.random.uniform(low=0,high=.5,size=100)
         y=np.linspace(1,20,100)+ np.random.uniform(low=0,high=1,size=100)
         print ('x = ',x)
         print ('y= ',y)

```

```

x = [ 1.34683976  1.12176759  1.51512398  1.55233174  1.40619168
      1.65075498  1.79399331  1.80243817  1.89844195  2.00100023
      2.3344038   2.22424872  2.24914511  2.36268477  2.49808849
      2.8212704   2.68452475  2.68229427  3.09511169  2.95703884
      3.09047742  3.2544361   3.41541904  3.40886375  3.50672677
      3.74960644  3.64861355  3.7721462   3.56368566  4.01092701
      4.15630694  4.06088549  4.02517179  4.25169402  4.15897504
      4.26835333  4.32520644  4.48563164  4.78490721  4.84614839
      4.96698768  5.18754259  5.29582013  5.32097781  5.0674106
      5.47601124  5.46852704  5.64537452  5.49642807  5.89755027
      5.68548923  5.76276141  5.94613234  6.18135713  5.96522091
      6.0275473   6.54290191  6.4991329   6.74003765  6.81809807
      6.50611821  6.91538752  7.01250925  6.89905417  7.31314433
      7.20472297  7.1043621   7.48199528  7.58957227  7.61744354
      7.6991707   7.85436822  8.03510784  7.80787781  8.22410224
      7.99366248  8.40581097  8.28913792  8.45971515  8.54227144
      8.6906456   8.61856507  8.83489887  8.66309658  8.94837987
      9.20890222  8.9614749   8.92608294  9.13231416  9.55889896
      9.61488451  9.54252979  9.42015491  9.90952569 10.00659591
      10.02504265 10.07330937  9.93489915 10.0892334  10.36509991]
y= [ 1.6635012   2.0214592   2.10816052  2.26016496  1.96287558

```

2.9554635	3.02881887	3.33565296	2.75465779	3.4250107
3.39670148	3.39377767	3.78503343	4.38293049	4.32963586
4.03925039	4.73691868	4.30098399	4.8416329	4.78175957
4.99765787	5.31746817	5.76844671	5.93723749	5.72811642
6.70973615	6.68143367	6.57482731	7.17737603	7.54863252
7.30221419	7.3202573	7.78023884	7.91133365	8.2765417
8.69203281	8.78219865	8.45897546	8.89094715	8.81719921
8.87106971	9.66192562	9.4020625	9.85990783	9.60359778
10.07386266	10.6957995	10.66721916	11.18256285	10.57431836
11.46744716	10.94398916	11.26445259	12.09754828	12.11988037
12.121557	12.17613693	12.43750193	13.00912372	12.86407194
13.24640866	12.76120085	13.11723062	14.07841099	14.19821707
14.27289001	14.30624942	14.63060835	14.2770918	15.0744923
14.45261619	15.11897313	15.2378667	15.27203124	15.32491892
16.01095271	15.71250558	16.29488506	16.70618934	16.56555394
16.42379457	17.18144744	17.13813976	17.69613625	17.37763019
17.90942839	17.90343733	18.01951169	18.35727914	18.16841269
18.61813748	18.66062754	18.81217983	19.44995194	19.7213867
19.71966726	19.78961904	19.64385088	20.69719809	20.07974319]

Find the expected value of x and the expected value of y.

#### Solution

Resource: (<https://revisionmaths.com/advanced-level-maths-revision/statistics/uniform-distribution>)

```
In [30]: # print(np.mean(x))
         print("x Expected Value: ", np.average(x))
         print("y Expected Value: ", np.average(y))
```

x Expected Value: 5.78253254159

y Expected Value: 11.0129816833

Find the variance for x and y.

```
In [31]: print("x Variance: ", np.var(x))
         print("y Variance: ", np.var(y))
```

x Variance: 7.03332752948

y Variance: 30.1139035755

Find the co-variance of x and y.

```
In [32]: print("x Co-Variance: ", np.cov(x))
         print("y Co-Variance: ", np.cov(y))
```

x Co-Variance: 7.104371241894797

y Co-Variance: 30.418084419706705

Assume that the number of dollars spent on car fuel is only linearly dependent on the miles driven. Write code that uses a linear predictor to calculate a predicted value of  $y$  for each  $x$ .

i.e.  $y_{predicted} = f(x) = mx + b$ .

In [ ]:

Predict  $y$  for each value in  $x$ , put the error into an array called  $y_{error}$ .

In [33]: `y_error =`

```
File "<ipython-input-33-a84fe419882b>", line 1
y_error =
^
```

SyntaxError: invalid syntax

Write code that calculates the root mean square error (RMSE).

In [34]: `rmse =`

```
File "<ipython-input-34-5059e2ec92a6>", line 1
rmse =
^
```

SyntaxError: invalid syntax

## 1.4 Pandas

### 1.4.1 Reading a File

Read in a CSV file called 'data3.csv' into a dataframe called `df`.

Data description \* Data source: <http://www.fao.org/nr/water/aquastat/data/query/index.html>  
\* Data, units \* GDP, current USD (CPI adjusted) \* NRI, mm/yr \* Population density, inhab/km<sup>2</sup>  
\* Total area of the country, 1000 ha = 10km<sup>2</sup> \* Total Population, unit 1000 inhabitants  
Display the first 10 lines of the dataframe.

In [52]: `import pandas as pd`

In [53]: `df = pd.read_csv('data3.csv')`  
`df[:10]`

```
Out[53]:
```

	Area	Area Id	Variable Name	Variable Id	Year	\
0	Argentina	9.0	Total area of the country	4100.0	1962.0	
1	Argentina	9.0	Total area of the country	4100.0	1967.0	
2	Argentina	9.0	Total area of the country	4100.0	1972.0	
3	Argentina	9.0	Total area of the country	4100.0	1977.0	

4	Argentina	9.0	Total area of the country	4100.0	1982.0
5	Argentina	9.0	Total area of the country	4100.0	1987.0
6	Argentina	9.0	Total area of the country	4100.0	1992.0
7	Argentina	9.0	Total area of the country	4100.0	1997.0
8	Argentina	9.0	Total area of the country	4100.0	2002.0
9	Argentina	9.0	Total area of the country	4100.0	2007.0

	Value	Symbol	Md
0	278040.0	E	NaN
1	278040.0	E	NaN
2	278040.0	E	NaN
3	278040.0	E	NaN
4	278040.0	E	NaN
5	278040.0	E	NaN
6	278040.0	E	NaN
7	278040.0	E	NaN
8	278040.0	E	NaN
9	278040.0	E	NaN

Display the column names.

```
In [54]: # FOR BETTER DISPLAY
         for name in df.columns:
             print(name)
```

```
Area
Area Id
Variable Name
Variable Id
Year
Value
Symbol
Md
```

### 1.4.2 Data Preprocessing

Create a mask of NAN values (i.e. apply `.isnull` on the dataframe). Inspect the mask for 'True' values, they denote NANs.

*Hint: You will notice that the last 8 rows and the last column ('Other') have NAN values. You can also use `df.tail()` to see the last row.*

Remove the bottom 8 rows from the dataframe because they contain NAN values. Also remove the column 'Other'.

```
In [55]: mask = pd.isnull(df)

df = df[:(-8)]
df = df.drop(['Md'], axis = 1)

df.tail()
```



```
Out[55]:
```

	Area	Area Id	Variable Name	\
385	United States of America	231.0	National Rainfall Index (NRI)	
386	United States of America	231.0	National Rainfall Index (NRI)	
387	United States of America	231.0	National Rainfall Index (NRI)	
388	United States of America	231.0	National Rainfall Index (NRI)	
389	United States of America	231.0	National Rainfall Index (NRI)	

	Variable Id	Year	Value	Symbol
385	4472.0	1981.0	949.2	E
386	4472.0	1984.0	974.6	E
387	4472.0	1992.0	1020.0	E
388	4472.0	1996.0	1005.0	E
389	4472.0	2002.0	938.7	E

All the columns in our dataframe are not required for analysis. Drop these columns: Area Id, Variable Id, and Symbol and save the new dataframe as df1.

```
In [56]: df.drop(['Area Id', 'Variable Id', 'Symbol'], axis = 1, inplace = True)
df1 = df
df1
```

```
Out[56]:
```

	Area	Variable Name	Year	\
0	Argentina	Total area of the country	1962.0	
1	Argentina	Total area of the country	1967.0	
2	Argentina	Total area of the country	1972.0	
3	Argentina	Total area of the country	1977.0	
4	Argentina	Total area of the country	1982.0	
5	Argentina	Total area of the country	1987.0	
6	Argentina	Total area of the country	1992.0	
7	Argentina	Total area of the country	1997.0	
8	Argentina	Total area of the country	2002.0	
9	Argentina	Total area of the country	2007.0	
10	Argentina	Total area of the country	2012.0	
11	Argentina	Total area of the country	2014.0	
12	Argentina	Total population	1962.0	
13	Argentina	Total population	1967.0	
14	Argentina	Total population	1972.0	
15	Argentina	Total population	1977.0	
16	Argentina	Total population	1982.0	
17	Argentina	Total population	1987.0	
18	Argentina	Total population	1992.0	
19	Argentina	Total population	1997.0	
20	Argentina	Total population	2002.0	
21	Argentina	Total population	2007.0	
22	Argentina	Total population	2012.0	
23	Argentina	Total population	2015.0	
24	Argentina	Population density	1962.0	
25	Argentina	Population density	1967.0	

26	Argentina	Population density	1972.0
27	Argentina	Population density	1977.0
28	Argentina	Population density	1982.0
29	Argentina	Population density	1987.0
..	...	...	...
360	United States of America	Population density	1972.0
361	United States of America	Population density	1977.0
362	United States of America	Population density	1982.0
363	United States of America	Population density	1987.0
364	United States of America	Population density	1992.0
365	United States of America	Population density	1997.0
366	United States of America	Population density	2002.0
367	United States of America	Population density	2007.0
368	United States of America	Population density	2012.0
369	United States of America	Population density	2015.0
370	United States of America	Gross Domestic Product (GDP)	1962.0
371	United States of America	Gross Domestic Product (GDP)	1967.0
372	United States of America	Gross Domestic Product (GDP)	1972.0
373	United States of America	Gross Domestic Product (GDP)	1977.0
374	United States of America	Gross Domestic Product (GDP)	1982.0
375	United States of America	Gross Domestic Product (GDP)	1987.0
376	United States of America	Gross Domestic Product (GDP)	1992.0
377	United States of America	Gross Domestic Product (GDP)	1997.0
378	United States of America	Gross Domestic Product (GDP)	2002.0
379	United States of America	Gross Domestic Product (GDP)	2007.0
380	United States of America	Gross Domestic Product (GDP)	2012.0
381	United States of America	Gross Domestic Product (GDP)	2015.0
382	United States of America	National Rainfall Index (NRI)	1965.0
383	United States of America	National Rainfall Index (NRI)	1969.0
384	United States of America	National Rainfall Index (NRI)	1974.0
385	United States of America	National Rainfall Index (NRI)	1981.0
386	United States of America	National Rainfall Index (NRI)	1984.0
387	United States of America	National Rainfall Index (NRI)	1992.0
388	United States of America	National Rainfall Index (NRI)	1996.0
389	United States of America	National Rainfall Index (NRI)	2002.0

	Value
0	2.780400e+05
1	2.780400e+05
2	2.780400e+05
3	2.780400e+05
4	2.780400e+05
5	2.780400e+05
6	2.780400e+05
7	2.780400e+05
8	2.780400e+05
9	2.780400e+05
10	2.780400e+05

11	2.780400e+05
12	2.128800e+04
13	2.293200e+04
14	2.478300e+04
15	2.687900e+04
16	2.899400e+04
17	3.132600e+04
18	3.365500e+04
19	3.583400e+04
20	3.788900e+04
21	3.997000e+04
22	4.209500e+04
23	4.341700e+04
24	7.656000e+00
25	8.248000e+00
26	8.913000e+00
27	9.667000e+00
28	1.043000e+01
29	1.127000e+01
..	...
360	2.214000e+01
361	2.317000e+01
362	2.430000e+01
363	2.549000e+01
364	2.678000e+01
365	2.834000e+01
366	2.995000e+01
367	3.132000e+01
368	3.202000e+01
369	3.273000e+01
370	6.050000e+11
371	8.620000e+11
372	1.280000e+12
373	2.090000e+12
374	3.340000e+12
375	4.870000e+12
376	6.540000e+12
377	8.610000e+12
378	1.100000e+13
379	1.450000e+13
380	1.620000e+13
381	1.790000e+13
382	9.285000e+02
383	9.522000e+02
384	1.008000e+03
385	9.492000e+02
386	9.746000e+02
387	1.020000e+03

```
388 1.005000e+03
389 9.387000e+02
```

```
[390 rows x 4 columns]
```

Display all the unique values in your new dataframe for these columns: Area, Variable Name, and Year.

Note the Countries and the Metrics (ie.recorded variables) represented in your dataset. *Hint: Use .unique( ) method.*

```
In [57]: print('The unique values in my new dataframe for Area is: \n', np.unique(df1[['Area']]))
         print('The unique values in my new dataframe for Variable Name is: \n', np.unique(df1[['Variable Name']]))
         print('The unique values in my new dataframe for Year is: \n', np.unique(df1[['Year']]))
```

The unique values in my new dataframe for Area is:

```
['Argentina' 'Australia' 'Germany' 'Iceland' 'Ireland' 'Sweden'
 'United States of America']
```

The unique values in my new dataframe for Variable Name is:

```
['Gross Domestic Product (GDP)' 'National Rainfall Index (NRI)'
 'Population density' 'Total area of the country' 'Total population']
```

The unique values in my new dataframe for Year is:

```
[1962. 1963. 1964. 1965. 1967. 1969. 1970. 1971. 1972. 1973.
 1974. 1975. 1977. 1978. 1979. 1981. 1982. 1983. 1984. 1985.
 1986. 1987. 1988. 1990. 1991. 1992. 1993. 1995. 1996. 1997.
 1998. 2000. 2001. 2002. 2007. 2012. 2014. 2015.]
```

Convert the Year column string values to pandas datetime objects, where only the year is specified.

*Hint: df1['Year'] = pd.to\_datetime(pd.Series(df1['Year']).astype(int),format='%Y').dt.year*

Run df1.tail() to see part of the result.

```
In [58]: df1['Year'] = pd.to_datetime(pd.Series(df1['Year']).astype(int),format='%Y').dt.year
         df1.tail()
```

```
Out[58]:
```

	Area	Variable Name	Year	Value
385	United States of America	National Rainfall Index (NRI)	1981	949.2
386	United States of America	National Rainfall Index (NRI)	1984	974.6
387	United States of America	National Rainfall Index (NRI)	1992	1020.0
388	United States of America	National Rainfall Index (NRI)	1996	1005.0
389	United States of America	National Rainfall Index (NRI)	2002	938.7

### 1.4.3 Extracting Statistics

Create a dataframe 'dftemp' to store rows where the Area is Iceland.

```
In [59]: dftemp = df1.query('Area=="Iceland"')
         dftemp
```

```

Out [59]:
      Area
166 Iceland      Total area of the country 1962 1.030000e+04
167 Iceland      Total area of the country 1967 1.030000e+04
168 Iceland      Total area of the country 1972 1.030000e+04
169 Iceland      Total area of the country 1977 1.030000e+04
170 Iceland      Total area of the country 1982 1.030000e+04
171 Iceland      Total area of the country 1987 1.030000e+04
172 Iceland      Total area of the country 1992 1.030000e+04
173 Iceland      Total area of the country 1997 1.030000e+04
174 Iceland      Total area of the country 2002 1.030000e+04
175 Iceland      Total area of the country 2007 1.030000e+04
176 Iceland      Total area of the country 2012 1.030000e+04
177 Iceland      Total area of the country 2014 1.030000e+04
178 Iceland      Total population          1962 1.826000e+02
179 Iceland      Total population          1967 1.974000e+02
180 Iceland      Total population          1972 2.099000e+02
181 Iceland      Total population          1977 2.221000e+02
182 Iceland      Total population          1982 2.331000e+02
183 Iceland      Total population          1987 2.469000e+02
184 Iceland      Total population          1992 2.599000e+02
185 Iceland      Total population          1997 2.728000e+02
186 Iceland      Total population          2002 2.869000e+02
187 Iceland      Total population          2007 3.054000e+02
188 Iceland      Total population          2012 3.234000e+02
189 Iceland      Total population          2015 3.294000e+02
190 Iceland      Population density        1962 1.773000e+00
191 Iceland      Population density        1967 1.917000e+00
192 Iceland      Population density        1972 2.038000e+00
193 Iceland      Population density        1977 2.156000e+00
194 Iceland      Population density        1982 2.263000e+00
195 Iceland      Population density        1987 2.397000e+00
196 Iceland      Population density        1992 2.523000e+00
197 Iceland      Population density        1997 2.649000e+00
198 Iceland      Population density        2002 2.785000e+00
199 Iceland      Population density        2007 2.965000e+00
200 Iceland      Population density        2012 3.140000e+00
201 Iceland      Population density        2015 3.198000e+00
202 Iceland      Gross Domestic Product (GDP) 1962 2.849165e+08
203 Iceland      Gross Domestic Product (GDP) 1967 6.212260e+08
204 Iceland      Gross Domestic Product (GDP) 1972 8.465069e+08
205 Iceland      Gross Domestic Product (GDP) 1977 2.226539e+09
206 Iceland      Gross Domestic Product (GDP) 1982 3.232804e+09
207 Iceland      Gross Domestic Product (GDP) 1987 5.565384e+09
208 Iceland      Gross Domestic Product (GDP) 1992 7.138788e+09
209 Iceland      Gross Domestic Product (GDP) 1997 7.596126e+09
210 Iceland      Gross Domestic Product (GDP) 2002 9.161798e+09
211 Iceland      Gross Domestic Product (GDP) 2007 2.129384e+10
212 Iceland      Gross Domestic Product (GDP) 2012 1.419452e+10

```

213	Iceland	Gross Domestic Product (GDP)	2015	1.659849e+10
214	Iceland	National Rainfall Index (NRI)	1967	8.160000e+02
215	Iceland	National Rainfall Index (NRI)	1971	9.632000e+02
216	Iceland	National Rainfall Index (NRI)	1975	1.010000e+03
217	Iceland	National Rainfall Index (NRI)	1981	9.326000e+02
218	Iceland	National Rainfall Index (NRI)	1986	9.685000e+02
219	Iceland	National Rainfall Index (NRI)	1991	1.095000e+03
220	Iceland	National Rainfall Index (NRI)	1997	9.932000e+02
221	Iceland	National Rainfall Index (NRI)	1998	9.234000e+02

Print the years when the National Rainfall Index (NRI) was > 950 or < 900 in Iceland using the dataframe you created in the previous question.

```
In [85]: print(dftemp[(dftemp["Variable Name"] == "National Rainfall Index (NRI)") & ((dftemp["V
```

	Area	Variable Name	Year	Value
214	Iceland	National Rainfall Index (NRI)	1967	816.0
215	Iceland	National Rainfall Index (NRI)	1971	963.2
216	Iceland	National Rainfall Index (NRI)	1975	1010.0
218	Iceland	National Rainfall Index (NRI)	1986	968.5
219	Iceland	National Rainfall Index (NRI)	1991	1095.0
220	Iceland	National Rainfall Index (NRI)	1997	993.2

Get all the rows of df1 (from the preprocessed data section of this notebook) where the Area is United States of America and store that into a new dataframe called df\_usa. Set the indices of the this dataframe to be the Year column.

*Hint: Use .set\_index()*

```
In [83]: df_usa = df1.query('Area=="United States of America").set_index('Year')
df_usa
```

```
Out[83]:
```

	Area	Variable Name	Value
Year			
1962	United States of America	Total area of the country	9.629090e+05
1967	United States of America	Total area of the country	9.629090e+05
1972	United States of America	Total area of the country	9.629090e+05
1977	United States of America	Total area of the country	9.629090e+05
1982	United States of America	Total area of the country	9.629090e+05
1987	United States of America	Total area of the country	9.629090e+05
1992	United States of America	Total area of the country	9.629090e+05
1997	United States of America	Total area of the country	9.629090e+05
2002	United States of America	Total area of the country	9.632030e+05
2007	United States of America	Total area of the country	9.632030e+05
2012	United States of America	Total area of the country	9.831510e+05
2014	United States of America	Total area of the country	9.831510e+05
1962	United States of America	Total population	1.918610e+05
1967	United States of America	Total population	2.037130e+05
1972	United States of America	Total population	2.132200e+05

1977	United States of America	Total population	2.230910e+05
1982	United States of America	Total population	2.339540e+05
1987	United States of America	Total population	2.454250e+05
1992	United States of America	Total population	2.579080e+05
1997	United States of America	Total population	2.728830e+05
2002	United States of America	Total population	2.884710e+05
2007	United States of America	Total population	3.016560e+05
2012	United States of America	Total population	3.147990e+05
2015	United States of America	Total population	3.217740e+05
1962	United States of America	Population density	1.993000e+01
1967	United States of America	Population density	2.116000e+01
1972	United States of America	Population density	2.214000e+01
1977	United States of America	Population density	2.317000e+01
1982	United States of America	Population density	2.430000e+01
1987	United States of America	Population density	2.549000e+01
1992	United States of America	Population density	2.678000e+01
1997	United States of America	Population density	2.834000e+01
2002	United States of America	Population density	2.995000e+01
2007	United States of America	Population density	3.132000e+01
2012	United States of America	Population density	3.202000e+01
2015	United States of America	Population density	3.273000e+01
1962	United States of America	Gross Domestic Product (GDP)	6.050000e+11
1967	United States of America	Gross Domestic Product (GDP)	8.620000e+11
1972	United States of America	Gross Domestic Product (GDP)	1.280000e+12
1977	United States of America	Gross Domestic Product (GDP)	2.090000e+12
1982	United States of America	Gross Domestic Product (GDP)	3.340000e+12
1987	United States of America	Gross Domestic Product (GDP)	4.870000e+12
1992	United States of America	Gross Domestic Product (GDP)	6.540000e+12
1997	United States of America	Gross Domestic Product (GDP)	8.610000e+12
2002	United States of America	Gross Domestic Product (GDP)	1.100000e+13
2007	United States of America	Gross Domestic Product (GDP)	1.450000e+13
2012	United States of America	Gross Domestic Product (GDP)	1.620000e+13
2015	United States of America	Gross Domestic Product (GDP)	1.790000e+13
1965	United States of America	National Rainfall Index (NRI)	9.285000e+02
1969	United States of America	National Rainfall Index (NRI)	9.522000e+02
1974	United States of America	National Rainfall Index (NRI)	1.008000e+03
1981	United States of America	National Rainfall Index (NRI)	9.492000e+02
1984	United States of America	National Rainfall Index (NRI)	9.746000e+02
1992	United States of America	National Rainfall Index (NRI)	1.020000e+03
1996	United States of America	National Rainfall Index (NRI)	1.005000e+03
2002	United States of America	National Rainfall Index (NRI)	9.387000e+02

Pivot the dataframe so that the unique Variable Name entries become the column entries. The dataframe values should be the ones in the Value column. Do this by running the lines of code below.

```
In [61]: df_usa=df_usa.pivot(columns='Variable Name',values='Value')
df_usa.head()
```

```
Out[61]: Variable Name  Gross Domestic Product (GDP)  National Rainfall Index (NRI)  \
Year
1962                6.050000e+11                NaN
1965                NaN                928.5
1967                8.620000e+11                NaN
1969                NaN                952.2
1972                1.280000e+12                NaN
```

```
Variable Name  Population density  Total area of the country  Total population
Year
1962                19.93                962909.0                191861.0
1965                NaN                NaN                NaN
1967                21.16                962909.0                203713.0
1969                NaN                NaN                NaN
1972                22.14                962909.0                213220.0
```

Rename the corresponding columns to ['GDP','NRI','PD','Area','Population'].

```
In [62]: df_usa = df_usa.rename(columns={"Gross Domestic Product (GDP)": "GDP",
                                         "National Rainfall Index (NRI)": "NRI",
                                         "Population density": "PD",
                                         "Total area of the country": "Area",
                                         "Total population": "Population"})
```

```
df_usa
```

```
Out[62]: Variable Name      GDP      NRI      PD      Area  Population
Year
1962      6.050000e+11      NaN  19.93  962909.0    191861.0
1965      NaN      928.5      NaN      NaN      NaN
1967      8.620000e+11      NaN  21.16  962909.0    203713.0
1969      NaN      952.2      NaN      NaN      NaN
1972      1.280000e+12      NaN  22.14  962909.0    213220.0
1974      NaN     1008.0      NaN      NaN      NaN
1977      2.090000e+12      NaN  23.17  962909.0    223091.0
1981      NaN      949.2      NaN      NaN      NaN
1982      3.340000e+12      NaN  24.30  962909.0    233954.0
1984      NaN      974.6      NaN      NaN      NaN
1987      4.870000e+12      NaN  25.49  962909.0    245425.0
1992      6.540000e+12    1020.0  26.78  962909.0    257908.0
1996      NaN     1005.0      NaN      NaN      NaN
1997      8.610000e+12      NaN  28.34  962909.0    272883.0
2002      1.100000e+13     938.7  29.95  963203.0    288471.0
2007      1.450000e+13      NaN  31.32  963203.0    301656.0
2012      1.620000e+13      NaN  32.02  983151.0    314799.0
2014      NaN      NaN      NaN  983151.0      NaN
2015      1.790000e+13      NaN  32.73      NaN    321774.0
```

Print the output of `df_usa.isnull().sum()`. This gives us the number of NAN values in each column. Replace the NAN values by 0, using `df_usa=df_usa.fillna(0)`. Print the output of `df_usa.isnull().sum()` again.



```
In [63]: print("Number of NAN values before: ", df_usa.isnull().sum())
         df_usa=df_usa.fillna(0)
         print("Number of NAN values after: ", df_usa.isnull().sum())
```

```
Number of NAN values before:  Variable Name
GDP          7
NRI          11
PD           7
Area         7
Population   7
dtype: int64
Number of NAN values after:  Variable Name
GDP          0
NRI          0
PD           0
Area         0
Population   0
dtype: int64
```

Calculate and print all the column averages and the column standard deviations.

```
In [64]: print("--- MEAN ---")
         for col in list(df_usa.columns.values):
             print(col, " : ", df_usa[col].mean())
         print()

         print("--- STANDARD DEVIATIONS ---")
         for col in list(df_usa.columns.values):
             print(col, " : ", df_usa[col].std())
```

```
--- MEAN ---
GDP : 4620894736842.105
NRI : 409.2736842105263
PD : 16.70157894736842
Area : 610314.7368421053
Population : 161513.42105263157
```

```
--- STANDARD DEVIATIONS ---
GDP : 6088655543027.329
NRI : 493.55150338260347
PD : 13.554620476328724
Area : 478948.1688578794
Population : 131380.53815298682
```

Using the df\_usa dataframe, multiply the Area by 10 (so instead of 1000 ha, the unit becomes 100 ha = 1km<sup>2</sup>). Store the result in place.

```
In [65]: df_usa.loc[:, 'Area'] *= 10
df_usa.tail()
```

```
Out[65]: Variable Name      GDP      NRI      PD      Area  Population
Year
2002      1.100000e+13  938.7  29.95  9632030.0    288471.0
2007      1.450000e+13    0.0  31.32  9632030.0    301656.0
2012      1.620000e+13    0.0  32.02  9831510.0    314799.0
2014      0.000000e+00    0.0   0.00  9831510.0         0.0
2015      1.790000e+13    0.0  32.73     0.0    321774.0
```

Create a new column in df\_usa called GDP/capita and populate it with the calculated GDP per capita. Round the results to two decimal points. Store the result in place.

```
In [69]: df_usa['GDP/capita'] = round(df_usa['GDP']/df_usa['Population']/1000, 2)
df_usa.head()
```

```
Out[69]: Variable Name      GDP      NRI      PD      Area  Population  GDP/capita
Year
1962      6.050000e+11    0.0  19.93  9629090.0    191861.0    3153.32
1965      0.000000e+00  928.5   0.00     0.0         0.0         NaN
1967      8.620000e+11    0.0  21.16  9629090.0    203713.0    4231.44
1969      0.000000e+00  952.2   0.00     0.0         0.0         NaN
1972      1.280000e+12    0.0  22.14  9629090.0    213220.0    6003.19
```

Create a new column in df\_usa called PD2 (i.e. population density 2). Calculate the population density. **Note: the units should be inhab/km<sup>2</sup>**. Round the results to two decimal point. Store the result in place.

```
In [72]: df_usa['PD2'] = round(df_usa['Population']*1000/df_usa['Area'], 2)
df_usa.head()
```

```
Out[72]: Variable Name      GDP      NRI      PD      Area  Population  GDP/capita  \
Year
1962      6.050000e+11    0.0  19.93  9629090.0    191861.0    3153.32
1965      0.000000e+00  928.5   0.00     0.0         0.0         NaN
1967      8.620000e+11    0.0  21.16  9629090.0    203713.0    4231.44
1969      0.000000e+00  952.2   0.00     0.0         0.0         NaN
1972      1.280000e+12    0.0  22.14  9629090.0    213220.0    6003.19

Variable Name      PD2
Year
1962      19.93
1965      NaN
1967      21.16
1969      NaN
1972      22.14
```

Find the maximum value and minimum value of the 'NRI' column in the USA (using pandas methods). What years do the min and max values occur in?

```
In [75]: print("MAX: ", df_usa["NRI"].max())
         print("MIN: ", df_usa[df_usa["NRI"] > 0]["NRI"].min())
         #print(list(df_usa.columns.values))
```

MAX: 1020.0

MIN: 928.5

## 1.5 Matplotlib

Create a dataframe called `icecream` that has column `Flavor` with entries `Strawberry`, `Vanilla`, and `Chocolate` and another column with `Price` with entries 3.50, 3.00, and 4.25.

```
In [76]: df = pd.DataFrame(data={'Flavor': ['Strawberry', 'Vanilla', 'Chocolate'],
                                'Price': [3.50, 3.00, 4.25]})
df
```

```
Out[76]:
```


	Flavor	Price
0	Strawberry	3.50
1	Vanilla	3.00
2	Chocolate	4.25

Create a bar chart representing the three flavors and their associated prices.

```
In [78]: objects = (df['Flavor'])
         price = list(df['Price'])

         plt.bar(np.arange(len(objects)), price, align='center', alpha=0.5)
         plt.xticks(np.arange(len(objects)), objects)
         plt.ylabel('Price')
         plt.title('Price of Flavors')

         plt.show()
```



Create 9 random plots. The top three should be scatter plots (one with green dots, one with purple crosses, and one with blue triangles). The middle three graphs should be a line graph, a horizontal bar chart, and a histogram. The bottom three graphs should be trigonometric functions (one sin, one cosine, one tangent).

```
In [117]: x = np.random.rand(50,1)
          y = np.random.rand(50,1)

          x_p = [1,2,4,6,8,10,12,14,16,19,20, 22, 24]
          y_p = [3,1,5,7,8,7,3,5,7,8,7,6,3]

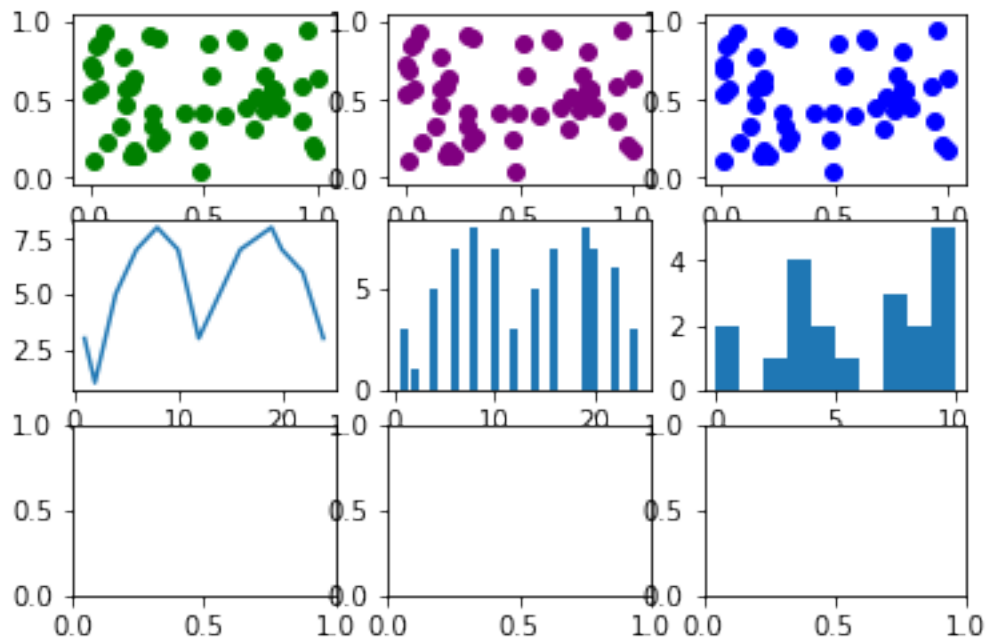
          f, ax = plt.subplots(nrows=3,ncols=3)

          N=20
          vals = np.random.randint(0,11,N)

          ax[0,0].scatter(x,y, c="Green")
          ax[0,1].scatter(x, y, c="Purple")
          ax[0,2].scatter(x, y, c="Blue")

          ax[1,0].plot(x_p,y_p)
          ax[1,1].bar(x_p, y_p)
          ax[1,2].hist(vals)
```

```
plt.show()
```



## 1.6 Extra Credit

Run the cell below to read in the data. See: <https://www.quantshare.com/sa-43-10-ways-to-download-historical-stock-quotes-data-for-free>

```
In [79]: df_google = pd.read_csv('https://finance.google.com/finance/historical?output=csv&q=goog')
df_apple = pd.read_csv('https://finance.google.com/finance/historical?output=csv&q=aapl')

df_disney = pd.read_csv('https://finance.google.com/finance/historical?output=csv&q=disney')
df_nike = pd.read_csv('https://finance.google.com/finance/historical?output=csv&q=nke')

df_apple.head()
```

```
Out[79]:
```

	Date	Open	High	Low	Close	Volume
0	8-Feb-18	160.29	161.00	155.03	155.15	53948375
1	7-Feb-18	163.08	163.40	159.07	159.54	51608580
2	6-Feb-18	154.83	163.72	154.00	163.03	68243838
3	5-Feb-18	159.10	163.88	156.00	156.49	72738522
4	2-Feb-18	166.00	166.80	160.10	160.50	86593825

Show a 3 x 3 correlation matrix for Nike, Apple, and Disney stock prices for the month of July, 2017.

Hint: Convert Date to a pandas datetime object. Change the indices of all the dataframes to Date. Use Date indices to filter rows. Create a new dataframe that stores values of the Close column from each dataframe. Use the Close column of each company's stock data to find the correlation using df.corr().

```
In [80]: df_google['Date'] = pd.to_datetime(pd.Series(df_google['Date']).astype(str),
                                             format='%d-%b-%y')
df_apple['Date'] = pd.to_datetime(pd.Series(df_apple['Date']).astype(str),
                                  format='%d-%b-%y')
df_disney['Date'] = pd.to_datetime(pd.Series(df_disney['Date']).astype(str),
                                   format='%d-%b-%y')
df_nike['Date'] = pd.to_datetime(pd.Series(df_nike['Date']).astype(str),
                                 format='%d-%b-%y')

In [81]: corr_google_apple = pd.concat([df_google["Close"], df_apple["Close"]], axis=1)
print("Corr Google - Apple")
print(corr_google_apple.corr(), "\n")

corr_google_disney = pd.concat([df_google["Close"], df_disney["Close"]], axis=1)
print("Corr Google - Disney")
print(corr_google_disney.corr(), "\n")

corr_google_nike = pd.concat([df_google["Close"], df_nike["Close"]], axis=1)
print("Corr Google - Nike")
print(corr_google_nike.corr(), "\n")

corr_apple_disney = pd.concat([df_apple["Close"], df_disney["Close"]], axis=1)
print("Corr Apple - Disney")
print(corr_apple_disney.corr(), "\n")

corr_apple_nike = pd.concat([df_apple["Close"], df_nike["Close"]], axis=1)
print("Corr Apple - Nike")
print(corr_apple_nike.corr(), "\n")

corr_disney_nike = pd.concat([df_disney["Close"], df_nike["Close"]], axis=1)
print("Corr Disney - Nike")
print(corr_disney_nike.corr(), "\n")
```

```
Corr Google - Apple
      Close      Close
Close  1.000000  0.871188
Close  0.871188  1.000000
```

```
Corr Google - Disney
      Close      Close
Close  1.000000 -0.195806
Close -0.195806  1.000000
```

```

Corr Google - Nike
      Close      Close
Close  1.000000  0.583454
Close  0.583454  1.000000

```

```

Corr Apple - Disney
      Close      Close
Close  1.000000 -0.297168
Close -0.297168  1.000000

```

```

Corr Apple - Nike
      Close      Close
Close  1.000000  0.48434
Close  0.48434  1.000000

```

```

Corr Disney - Nike
      Close      Close
Close  1.000000  0.389781
Close  0.389781  1.000000

```

Show the same correlation matrix but over different time periods. 1. the last 20 days  
2. the last 80 days

```
In [ ]: df_google["Date" > "10-Feb-18"]
```

Change the code so that it accepts a list of any stock symbols (i.e. ['NKE', 'APPL', 'DIS', ... ]) and creates a correlation matrix for the past 100 days.