



Data X

Introduction to NLP-II (Word2vec)

By Sana Iqbal

Review

- ❑ Goal of NLP:
Understand the meaning of the text → that is to get to the **semantic level analysis**
- ❑ We looked at **bag of words** model, which keeps the count of words in a document. Disadvantages:
Loss of the ordering of words --> **Ignore semantics** of the words

Information about order and context of words is important to understand the document.

Discrete Representation Of Words

- ❑ We represent words in our corpus as **atomic symbols** i.e each word is independant.

For Corpus : 1. 'I love knitting' 2. 'I love dogs'

- ❑ Vocabulary, $V = [i, \text{love}, \text{knitting}, \text{dogs}]$
- ❑ If we want to represent them as numbers in our machine we assign them an id. Eg:
I=1, love=2, knitting=3, dogs=4
- ❑ *But words have no ordinal relationship*

Vector Representation Of Words

❑ One-hot-Encode

❑ We can represent these **words as vectors**.

We can say in the vocabulary space $V = [i, \text{love}, \text{knitting}, \text{dogs}]$

i=	[1,0,0,0]
love=	[0,1,0,0]
knitting=	[0,0,1,0]
dogs=	[0,0,0,1]

Discrete Representation Of Text

- But **size of One-hot-Encoded** word vectors will depend on the corpus vocabulary size.
- We want you use NLP on large corpora that exist in the world, for example the *IMDB reviews*, *the yelp reviews*, *wikipedia articles* etc or **Google corpus**: a vocabulary of **3 million words** -Google News dataset

Problems in Discrete Representation Of Text

1. Vocabulary size is big.

We will end up having very very large sized **sparse vectors**.

- [illegible]


2. We are not able to capture any **semantic relations** between words.

To capture similarity between good & nice using the above vectors:

Cosine similarity= $\text{Dot}(\text{good}, \text{nice}) = 0$

Distributional Representation of Words

- ❑ We want to represent a word as **a vector** that **encodes its meaning**.
- ❑ Every word is represented as a function of other words.
- ❑ To do that we rely on **Distributional similarity** concept.
- ❑ The idea is that if we look at the different contexts in which a word appears or is used in a language, we will be able to infer its meaning.



Distributional Representation of Words cntd ...

Example:

Eating healthy is a key to fitness.

Junk **eating** causes obesity.

If you stop **eating**, you will die.

Too much **eating** will make you obese.

Not all cultures use spoons for **eating** food.

Eating seen in context of healthy, junk, food, fitness, spoons, die etc. gives the idea of its meaning.

Distributional Representation of Words cntd ...

JR Firth, a British linguist: "You shall know a word by the company it keeps."

- ❑ In **Distributional Representation**, we represent words as vectors that capture the *context* of these words in the corpus.
- ❑ **Basic Method:** We can count the number of times the word groups co-occur in the corpus. ----**CO-OCCURRENCE MATRIX**

Distributional Representation of Words --- Co-Occurrence Matrix

Corpus: I like deep learning. I like NLP. I enjoy flying.

Window size: 1

Vocab: I, like, enjoy, deep, learning, NLP, flying

Co-Occurrence Matrix


counts	I	like	enjoy	deep	learning	NLP	flying
I	0	2	1	0	0	0	0
like	2	0	0	1	0	1	0
enjoy	1	0	0	0	0	0	1
deep	0	1	0	0	1	0	0
learning	0	0	0	1	0	0	0
NLP	0	1	0	0	0	0	0
flying	0	0	1	0	0	0	0

→ Vector Representation of *like*

Issue:

- Again High dimensional vectors for a large corpus

Ref: <http://mysite.science.uottawa.ca/phofstra/MAT2342/SVDproblems.pdf>



Distributional Representation of Words --- Co-Occurrence Matrix

Corpus: *the quick brown fox jumped over the lazy dog and killed it*

Co-Occurrence Matrix

	the	quick	brown	fox	jumped	over	lazy	dog	and	killed	it
the	0	1	0	0	0	1	1	0	0	0	0
quick	1	0	1	0	0	0	0	0	0	0	0
brown	0	1	0	1	0	0	0	0	0	0	0
fox	0	0	1	0	1	0	0	0	0	0	0
jumped	0	0	0	1	0	1	0	0	0	0	0
over	1	0	0	0	1	0	1	0	0	0	0
lazy	1	0	0	0	0	1	0	1	0	0	0
dog	0	0	0	0	0	0	1	0	1	0	0
and	0	0	0	0	0	0	0	1	0	1	0
killed	0	0	0	0	0	0	0	0	1	0	1
it	0	0	0	0	0	0	0	0	0	1	0

→ Vector Representation of **the**

Issue:

- High dimensional vectors for a large corpus

WORD2VEC

By Mikolov et al.

- ❑ **WORD2VEC** is a method of creating distributional representations of words called **word embeddings**, using backpropagation.
- ❑ Eg: apple = $[0.2, 0.35, 0.1, -0.2, 0.15, 0.4]^T$

Review: A simple Feed Forward Neural Net


1. Accepts the Input.

2. Creates different combinations of **weighted input**

3. Induces **Non-linearity** in the input using Activations.

4. Uses backpropagation to update weights.

5. Uses the trained model to give Output on test data



Schematic of a multi-layer perceptron.

NOTE: *These are based on the combination of inputs*

WORD2VEC

- ❑ **Neural Net Models** that aim to **predict contextual words**/word of the input words/word.
- ❑ Two algorithms of word2vec:
 1. Skip-gram
 2. Continuous Bag of words (Cbow)

Skip-gram:

the task is "***predicting the context given a word***".

❑ $p(\text{context words} | \text{word})$

CBOW:

the task as "***predicting the word given its context***".

❑ $p(\text{word} | \text{context words})$

WORD2VEC

Input : word(s) **Output:** contextual word(s)

Like any other Neural Network model:

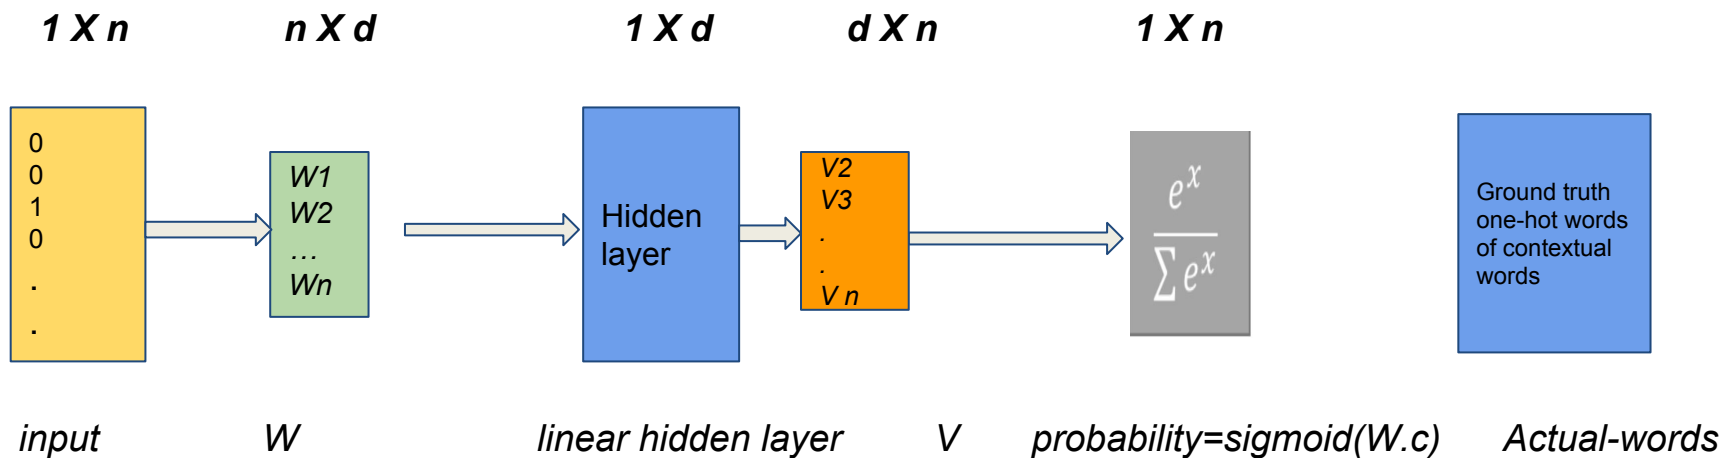
1. Word2vec Model is parametrized by weights.
2. Trained using a loss/ objective function
3. Weights are tuned to minimize loss BETWEEN actual contextual word and predicted contextual word.

SPECIAL:These parameter/weight matrix between input and hidden layer of the model is called '**word vector representations**' of words.



WORD2VEC - Skip Gram Model

- ❑ **Input:** One hot encoded word, c
- ❑ Weight Matrix: V and W
- ❑ Objective Function= maximize $p(\text{context word} | c)$
- ❑ $\text{probability}(x_i | c) = \text{softmax}(x_i \cdot c)$



One hot to Word2vec

One hot vectors

N -dimensional

Hidden Layer
Weight Matrix

D neurons

*Word Vector
Lookup Table!*




D -dimensional



N words

N words

N -words



Skip gram example data:

If we have:

n = Vocabulary in the corpus= 11

d = Word vector dimension=3

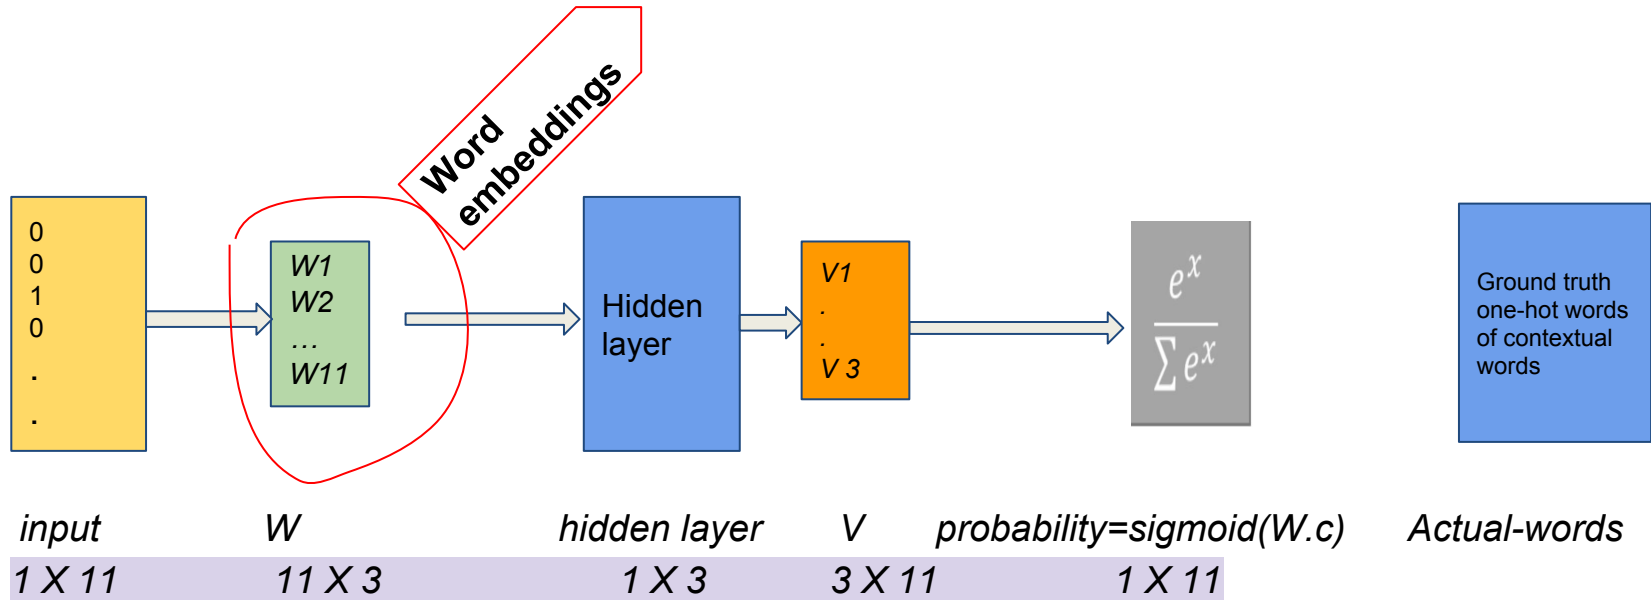
w= Window size on each side=1

Corpus: *the quick brown fox jumped over the lazy dog and killed it*




Output	Input
[the, brown]	quick
[quick, fox]	brown
[brown, jumped]	fox
...	


Skip gram example data:



Results with word2vec in the original paper trained on Google news dataset



From Mikolov *et al.*
(2013a)

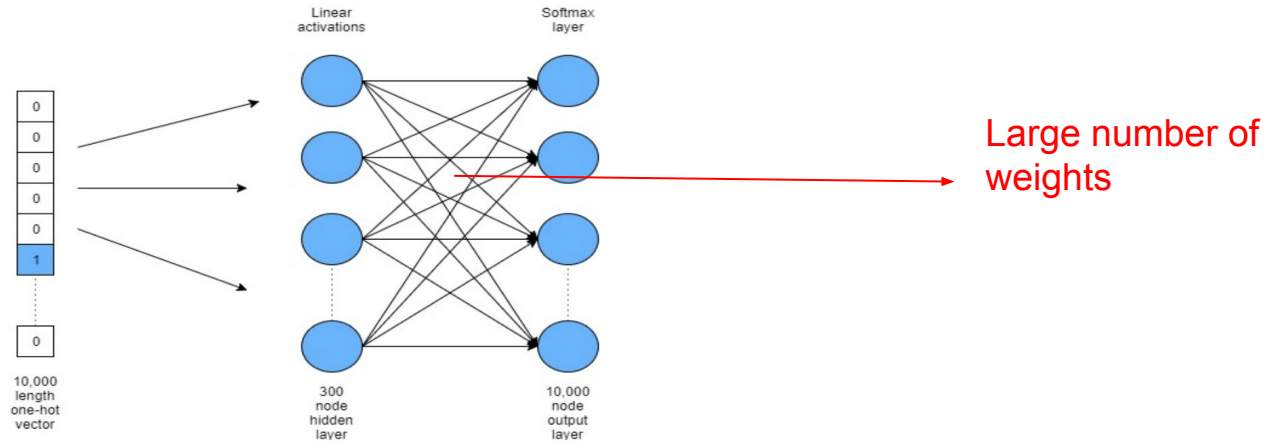


More..



Negative Sampling Introduction

- ❑ **Problem** with using a vanilla skip gram
Full softmax output layer ----> **computationally expensive.**
- ❑ When the output is a n-word one-hot vector, large number of weights that need to be updated in any gradient based training of the output layer.



Negative Sampling

Instead of constructing a network that outputs a multi-class softmax layer, **we change it into a simple binary classifier.**

- ❑ **Input:** [*a target word and a real or negative context word*]
- ❑ **Output:** [0 or 1 based on *a real or negative context word*]
- ❑ An embedding layer
- ❑ Similarity Operation: To train the model to assign similar words with similar embedding vectors.
- ❑ The output sigmoid layer [0,1]

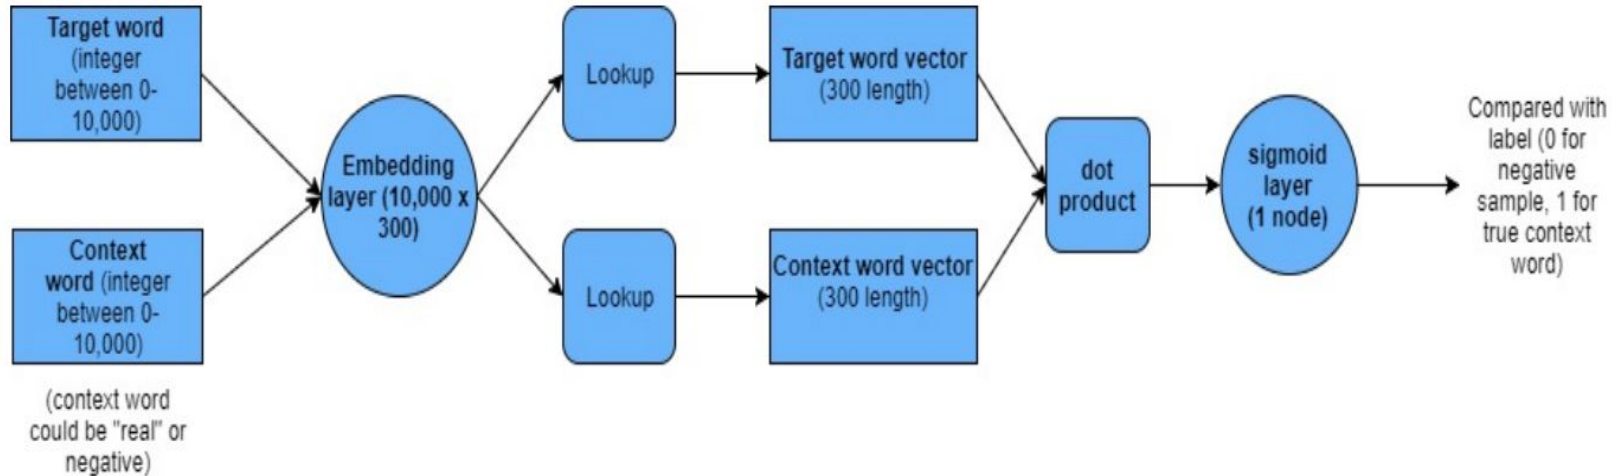
Negative Sampling

Corpus: *the quick brown fox jumped over the lazy dog and killed it*

Window size= 2 i.e 1 on each side of the input word

Input (target word, context)	Output (label)
[quick, brown]	1
[quick, dog]	0
[brown, fox]	1

Negative Sampling



Using pretrained word embeddings:

1. We can also use the word embeddings from pretrained models, eg the model trained on Google Data is available in many packages .
2. These are useful when we are working in the same domain or our own corpus is very small.

